

PROJECT DOCUMENTATION

AMAZON FOOD REVIEW SENTIMENT ANALYSIS USING NATURAL LANGUAGE PROCESSING

An Independent Self-Directed Project

Built and Documented by:

SHUBHRANSHU

GitHub: Shubhranshu331

Email: shubhranshu.331@gmail.com

2024 – 2025

ABOUT THIS PROJECT

This is a fully independent, self-directed project. It was conceived, designed, built, and documented entirely by me — without the guidance of any teacher, supervisor, mentor, or institution. No college, university, or organization is affiliated with this work in any capacity.

I undertook this project out of personal curiosity and a genuine desire to understand how Natural Language Processing and sentiment analysis work in practice on real-world data. Every line of code, every analytical decision, and every word of this documentation reflects my own research, thinking, and effort.

The tools, libraries, and datasets used are all freely available open-source resources. I am grateful to the open-source community — the developers of Python, pandas, NLTK, matplotlib, seaborn, and the WordCloud library — whose work made this project possible. I am also grateful to Julian McAuley and Jure Leskovec for making the Amazon Fine Food Reviews dataset publicly available on Kaggle, and to Hutto and Gilbert for open-sourcing the VADER sentiment tool.

I declare that all code, analysis, and documentation in this project are my own original work. Any external concepts or methods referenced are properly cited in the References section.

Shubhranshu

A PERSONAL NOTE

I want to be upfront about how this project came to exist. I am a self-learner. I did not follow a curriculum, sit in a classroom, or get grades for this. I built this project because I was genuinely curious about how machines can understand human emotions expressed in text — and specifically, whether a computer could look at a food review and correctly identify whether the person writing it was happy or unhappy with the product.

I started by reading about Natural Language Processing online, worked through documentation for the relevant Python libraries, and experimented with the code until I understood what was happening at each step. The Amazon Fine Food Reviews dataset appealed to me because it is large, real-world, and full of genuinely interesting variation — people write about food in deeply personal ways, and I wanted to see how a rule-based NLP tool would cope with that variation.

The process taught me far more than I expected — not just about sentiment analysis and NLP, but about data cleaning, exploratory analysis, visualization, and the importance of questioning your own results. The moments where VADER got things wrong were often the most interesting, because they revealed something real about the limits of automated language understanding.

I wrote this documentation to a professional standard not because anyone required it of me, but because I believe that good work deserves to be clearly explained. If you are reading this as a potential collaborator, employer, or fellow self-learner — I hope it gives you a clear and honest picture of what I built and what I learned.

Shubhranshu

ABSTRACT

Sentiment analysis, a cornerstone application of Natural Language Processing (NLP), enables automated identification of subjective information within textual data. In the era of digital commerce, online customer reviews represent an invaluable source of structured opinion data that individuals and organizations can leverage for product insights, market intelligence, and customer sentiment monitoring. This project presents a comprehensive sentiment analysis system applied to the Amazon Fine Food Reviews dataset — a publicly available corpus of approximately 568,454 consumer product reviews collected over a span of more than a decade. This project was built entirely independently, as a self-directed learning and portfolio exercise, with no institutional affiliation or academic supervision.

The primary analytical methodology employed in this project is the Valence Aware Dictionary and sEntiment Reasoner (VADER), a rule-based sentiment analysis tool specifically designed and optimized for social media and short consumer-text contexts. VADER scores each review text along four dimensions: negativity, neutrality, positivity, and a composite compound score, which is then used to classify each review as Positive, Neutral, or Negative according to empirically established threshold values (compound ≥ 0.05 for Positive; compound ≤ -0.05 for Negative; otherwise Neutral).

The project workflow encompasses several critical phases: data loading and profiling, rigorous data cleaning (removal of null entries, deduplication, and working sample generation), exploratory data analysis (star rating distribution, review length analysis, helpfulness ratio computation), text preprocessing for word-frequency analysis (lowercasing, HTML tag removal, punctuation stripping, and stopword removal), VADER sentiment scoring, and an extensive visualization suite including sentiment distribution bar charts, compound score density plots stratified by star rating, a cross-tabulated heatmap of star ratings versus VADER sentiment labels, word clouds per sentiment category, and top-word frequency bar plots.

Analysis of a representative 10,000-review sample demonstrates strong alignment between VADER-assigned sentiment labels and original star ratings: the agreement rate between VADER sentiment and star-rating-derived expected sentiment reached approximately 77% overall, with high-star (4-5 star) reviews predominantly classified as Positive and low-star (1-2 star) reviews predominantly classified as Negative. Edge cases of misclassification, such as sarcastic negative reviews receiving positive VADER scores, were identified and discussed, highlighting the known limitations of lexicon-based approaches in handling figurative language.

This project demonstrates that a lexicon-based NLP approach can achieve robust sentiment classification without requiring labeled training data, making it highly applicable in scenarios with limited annotated corpora. The findings, visualizations, and analytical insights produced in this study serve as a practical foundation for future research integrating supervised machine learning classifiers and deep learning architectures for further improvement in sentiment classification accuracy.

Keywords: Sentiment Analysis, Natural Language Processing, VADER, Amazon Fine Food Reviews, Text Classification, Opinion Mining, Lexicon-Based Analysis, Self-Directed Project, Python, Open-Source NLP.

TABLE OF CONTENTS

CHAPTER 1: INTRODUCTION	13
1.1 Overview of Machine Learning	13
1.2 Introduction to Natural Language Processing.....	14
1.3 Importance of Sentiment Analysis.....	15
1.4 Real-World Applications of Sentiment Analysis.....	16
1.5 Importance of Customer Reviews in E-Commerce	17
1.6 Need for Automated Review Classification	17
1.7 Challenges in Text Classification Systems	18
CHAPTER 2: PROBLEM STATEMENT.....	20
CHAPTER 3: OBJECTIVES AND SCOPE OF THE PROJECT	22
3.1 Objectives of the Project.....	22
Objective 1: Dataset Profiling and Quality Assessment	22
Objective 2: Text Preprocessing Pipeline Implementation.....	22
Objective 3: VADER Sentiment Scoring and Classification.....	22
Objective 4: Exploratory Data Analysis and Visualization	22
Objective 5: Sentiment-Rating Alignment Analysis.....	23
Objective 6: Critical Analysis of Results and Future Directions	23
3.2 Scope of the Project	23
CHAPTER 4: LITERATURE SURVEY.....	25
4.1 Introduction to the Literature Survey	25
4.2 Review of Related Work.....	25
4.2.1 Foundational Work: Pang, Lee, and Vaithyanathan (2002)	25
4.2.2 Turney (2002): Semantic Orientation via Mutual Information.....	25
4.2.3 Hutto and Gilbert (2014): VADER.....	26
4.2.4 McAuley and Leskovec (2013): Amazon Fine Food Reviews	26
4.2.5 Socher et al. (2013): Recursive Neural Networks for Sentiment.....	26
4.2.6 Devlin et al. (2019): BERT and Its Impact on NLP	27
4.2.7 Maas et al. (2011): Large Movie Review Dataset for SA.....	27
4.2.8 Bing Liu (2012): Comprehensive Survey of Sentiment Analysis	27
4.2.9 Rotten Tomatoes Multi-Aspect Sentiment (2020).....	28
4.2.10 Deep Learning Approaches: CNN and LSTM Models.....	28
4.3 Comparison Table of Related Works.....	28
CHAPTER 5: EXISTING SYSTEM AND PROPOSED SYSTEM.....	30
5.1 Existing System	30

5.1.1 Drawbacks of Existing Systems.....	30
5.2 Proposed System.....	31
CHAPTER 6: DATASET DESCRIPTION	32
6.1 Dataset Source and Background	32
6.2 Dataset Structure and Features.....	32
6.3 Score Distribution Analysis	33
6.4 Missing Value Analysis	34
CHAPTER 7: EXPLORATORY DATA ANALYSIS	35
7.1 Overview of EDA	35
7.2 Star Rating Distribution	35
7.3 Review Length Analysis	36
7.4 Helpfulness Ratio Analysis.....	37
CHAPTER 8: DATA CLEANING, PREPROCESSING AND NLP TECHNIQUES	38
8.1 Data Cleaning.....	38
8.1.1 Removal of Null Review Texts.....	38
8.1.2 Deduplication of Review Texts	38
8.1.3 Working Sample Generation.....	39
8.2 Text Preprocessing Pipeline.....	39
8.3 VADER Sentiment Analysis Framework	40
8.4 Text Vectorization and Feature Engineering Concepts	41
8.4.1 Bag of Words (BoW)	41
8.4.2 TF-IDF Vectorization	42
8.4.3 Word Embeddings	42
CHAPTER 9: SYSTEM ARCHITECTURE AND IMPLEMENTATION.....	44
9.1 Overall System Architecture.....	44
9.2 Implementation Details.....	44
9.2.1 Data Ingestion	44
9.2.2 Sampling Strategy	44
9.2.3 VADER Integration	45
9.2.4 Visualization Suite	45
CHAPTER 10: EXPERIMENTAL RESULTS AND ANALYSIS.....	46
10.1 Sentiment Distribution Results	46
10.2 Compound Score Distribution by Star Rating	47
10.3 Star Rating vs. VADER Sentiment Heatmap	47
10.4 Average Compound Score per Star Rating.....	48
10.5 Word Frequency and Word Cloud Analysis.....	49
10.6 VADER vs. Star Rating Agreement Analysis	51

10.7 Misclassification Analysis	51
CHAPTER 11: RESULT DISCUSSION, ADVANTAGES AND LIMITATIONS...	53
11.1 Result Discussion.....	53
11.2 Advantages of the Proposed System.....	54
11.3 Limitations of the Proposed System	54
CHAPTER 12: APPLICATIONS AND FUTURE SCOPE	56
12.1 Applications of Sentiment Analysis.....	56
12.2 Future Scope	57
CHAPTER 13: CONCLUSION	58
REFERENCES	60
APPENDIX A: CODE SNIPPETS AND SAMPLE OUTPUTS	62
A.1 Complete Import and Setup Code.....	62
A.2 Data Loading and Profiling Code	62
A.3 Text Preprocessing Function.....	62
A.4 VADER Sentiment Scoring Code.....	63
A.5 Agreement Rate Calculation	63
A.6 Sample Prediction Output	63

LIST OF FIGURES

Figure 1.1	Architecture of a Typical Sentiment Analysis Pipeline	10
Figure 6.1	Amazon Fine Food Reviews Dataset Schema	30
Figure 7.1	Distribution of Star Ratings in the Dataset	35
Figure 7.2	Distribution of Review Length (Character Count)	36
Figure 7.3	Average Review Length by Star Rating	37
Figure 8.1	Text Preprocessing Pipeline	41
Figure 9.1	VADER Compound Score Distribution by Star Rating	48
Figure 9.2	Overall Sentiment Distribution (VADER Labels)	49
Figure 9.3	Star Rating vs. VADER Sentiment Heatmap	50
Figure 9.4	Average VADER Compound Score per Star Rating	51
Figure 9.5	Word Cloud — Positive Reviews	52
Figure 9.6	Word Cloud — Negative Reviews	53
Figure 9.7	Word Cloud — Neutral Reviews	54
Figure 9.8	Top 15 Most Frequent Words per Sentiment Category	55
Figure 10.1	VADER vs. Star Rating Agreement Rate Analysis	57
Figure 10.2	Misclassification Examples — 1-Star Reviews Labelled Positive	59

LIST OF TABLES

Table 1.1	Comparison of Sentiment Analysis Approaches	8
Table 4.1	Summary of Literature Survey (Related Works)	23
Table 4.2	Comparative Analysis of Prior Research	26
Table 6.1	Dataset Column Descriptions	31
Table 6.2	Score Distribution in the Full Dataset	32
Table 6.3	Sample Review Records	33
Table 8.1	Text Preprocessing Steps and Their Rationale	40
Table 8.2	VADER Scoring Component Definitions	45
Table 9.1	VADER Sentiment Label Distribution in Sample	49
Table 9.2	Cross-Tabulation: Star Rating vs. VADER Sentiment	50
Table 10.1	VADER vs. Expected Sentiment Agreement by Star Rating	58
Table 10.2	Sample Misclassified Reviews	59
Table 11.1	Summary of Proposed System Performance	62

LIST OF ABBREVIATIONS

Abbreviation	Full Form
NLP	Natural Language Processing
VADER	Valence Aware Dictionary and sEntiment Reasoner
ML	Machine Learning
DL	Deep Learning
TF-IDF	Term Frequency – Inverse Document Frequency
BoW	Bag of Words
SVM	Support Vector Machine
LR	Logistic Regression
NB	Naive Bayes
RF	Random Forest
CSV	Comma-Separated Values
API	Application Programming Interface
EDA	Exploratory Data Analysis
HTML	HyperText Markup Language
NLTK	Natural Language Toolkit
ROC	Receiver Operating Characteristic
AUC	Area Under the Curve
F1	F1-Score (Harmonic Mean of Precision and Recall)
GPU	Graphics Processing Unit
AI	Artificial Intelligence

CHAPTER 1: INTRODUCTION

1.1 Overview of Machine Learning

Machine Learning (ML) is a transformative branch of artificial intelligence that empowers computational systems to learn from data and improve their performance on specific tasks without being explicitly programmed with domain-specific rules. The fundamental premise of machine learning is that algorithms can infer patterns, structures, and regularities from large datasets and subsequently generalize these learned representations to previously unseen data instances. This capability distinguishes machine learning fundamentally from traditional rule-based programming, where human experts must manually encode every decision criterion into the software system.

The history of machine learning as a formal discipline traces back to the foundational work of Alan Turing, Frank Rosenblatt, and Arthur Samuel in the mid-twentieth century. However, it was the convergence of three modern trends — the exponential growth of digital data, unprecedented advances in computational hardware, and the development of sophisticated optimization algorithms — that propelled machine learning into the center of technological innovation in the twenty-first century. Today, machine learning algorithms underpin an extraordinary range of applications, from computer vision and autonomous vehicles to drug discovery and financial forecasting.

Machine learning methodologies are broadly categorized into three learning paradigms. Supervised learning involves training a model on a labeled dataset, where each input instance is paired with a known output label, enabling the model to learn a mapping function from inputs to outputs. Prominent supervised algorithms include linear regression, logistic regression, support vector machines, decision trees, random forests, and deep neural networks. Unsupervised learning, by contrast, operates on unlabeled data and seeks to discover hidden structure, patterns, or clusters within the data without explicit output guidance. Clustering algorithms such as k-means, hierarchical clustering, and density-based methods are classic examples. Reinforcement learning constitutes a third paradigm, in which an agent learns to make sequential decisions by

interacting with an environment and receiving scalar reward signals that guide behavior optimization toward a long-term objective.

In the domain of text analytics and language understanding, machine learning has catalyzed remarkable breakthroughs. The field of Natural Language Processing has benefited enormously from both classical machine learning approaches and modern deep learning architectures, enabling computers to read, understand, and generate human language with increasingly human-like proficiency.

1.2 Introduction to Natural Language Processing

Natural Language Processing (NLP) is the interdisciplinary subfield of computer science, linguistics, and artificial intelligence concerned with enabling computers to understand, interpret, and generate human language in a meaningful and useful way. Natural language, unlike formal programming languages, is characterized by extraordinary richness, ambiguity, context-dependence, and variability. The same concept can be expressed through countless syntactic structures; the meaning of words shifts dramatically based on context; idioms, sarcasm, metaphors, and cultural references add further layers of complexity that challenge automated interpretation.

The scope of NLP encompasses a wide spectrum of tasks. At the lexical level, tokenization decomposes continuous text into discrete units (tokens), part-of-speech tagging assigns grammatical categories to each token, and morphological analysis handles word inflection and derivation. At the syntactic level, parsing constructs hierarchical grammatical representations of sentences. At the semantic level, named entity recognition identifies and classifies entities such as persons, organizations, and locations, while semantic role labeling determines the relationships between predicates and their arguments. At the discourse level, coreference resolution links pronouns and noun phrases to the entities they denote, and discourse parsing models the logical and rhetorical structure of multi-sentence texts.

Contemporary NLP is dominated by large-scale pre-trained language models, including BERT, GPT, RoBERTa, XLNet, and their successors. These models, trained on vast corpora of text using self-supervised objectives, acquire rich linguistic knowledge that

can be fine-tuned for downstream tasks with remarkably small amounts of labeled data. However, for many practical applications, especially those with domain-specific lexicons or operating under computational constraints, lexicon-based and classical machine learning NLP approaches remain highly relevant and competitive. The present project employs the VADER lexicon-based tool, which was specifically engineered and validated for sentiment classification in consumer and social media text, making it an appropriate choice for the Amazon food review domain.

1.3 Importance of Sentiment Analysis

Sentiment analysis, also referred to as opinion mining, is the computational process of identifying, extracting, and quantifying subjective information — such as opinions, emotions, attitudes, and evaluations — from textual sources. It represents one of the most practically impactful applications of NLP, enabling organizations to automatically process the vast volumes of user-generated text produced across social media platforms, e-commerce portals, news commentary sections, and customer support channels.

The commercial imperative for sentiment analysis has grown in direct proportion to the explosion of digital text data. Prior to the digital era, understanding public opinion required expensive and time-consuming surveys, focus groups, and market research studies. The proliferation of online platforms has rendered this traditional approach both insufficient in scale and inadequate in timeliness. Organizations now require near-real-time insights from millions of data points to remain competitive. Sentiment analysis provides the automated, scalable mechanism for converting unstructured text into actionable quantitative signals.

At its most fundamental level, sentiment analysis classifies text along a polarity dimension: positive, negative, or neutral. More sophisticated approaches extend this to fine-grained sentiment (e.g., a five-level scale from strongly negative to strongly positive), aspect-based sentiment analysis (which identifies sentiment expressed toward specific product attributes rather than the product as a whole), and emotion detection (which maps text to discrete emotional categories such as joy, sadness, anger, fear, and surprise). The granularity of analysis selected depends on the downstream application requirements and the availability of annotated training data.

1.4 Real-World Applications of Sentiment Analysis

The applications of sentiment analysis span virtually every industry sector. In e-commerce and retail, sentiment analysis of product reviews enables businesses to identify product defects, benchmark against competitors, and prioritize product improvement initiatives based on systematic analysis of customer feedback. Platforms such as Amazon, Flipkart, and Alibaba process millions of reviews daily, making automated sentiment analysis indispensable for quality management at scale.

In the financial services industry, sentiment analysis of news articles, earnings call transcripts, and social media posts has emerged as a significant component of algorithmic trading strategies. Research has consistently demonstrated that the sentiment expressed in financial news and analyst reports carries predictive information about short-term stock price movements, enabling quantitative hedge funds to generate alpha by incorporating textual sentiment signals into their models.

In public health and social policy, sentiment analysis of social media posts has been applied to monitor population mental health trends, detect early signals of disease outbreaks, and evaluate public attitudes toward vaccination programs and health interventions. During the COVID-19 pandemic, sentiment analysis was extensively employed to track public compliance with lockdown measures and assess the effectiveness of government communications.

In politics and governance, sentiment analysis of news media, parliamentary transcripts, and social media enables political analysts and strategists to monitor public opinion on policy issues, track shifts in political sentiment across demographic groups, and evaluate the impact of political events on electoral sentiment. Brand monitoring applications track the sentiment associated with corporate brands across media channels, enabling communications teams to respond proactively to reputational threats.

1.5 Importance of Customer Reviews in E-Commerce

Customer reviews occupy a uniquely important position in the digital commerce ecosystem. Unlike traditional retail environments, where consumers could physically inspect products before purchase and seek advice from knowledgeable sales personnel, online shopping inherently involves significant information asymmetry between sellers and buyers. Customer reviews serve as the primary mechanism for bridging this information gap, providing prospective buyers with peer assessments of product quality, performance, durability, and value.

Empirical research has consistently demonstrated the substantial economic impact of online reviews. Studies indicate that a one-star increase in a restaurant's average Yelp rating is associated with a five to nine percent increase in revenue. Similarly, product ratings on Amazon exert a significant influence on purchase conversion rates, with even fractional improvements in average star rating yielding measurable uplift in sales volume. The volume of reviews also carries independent informational value beyond average rating, as products with a larger number of reviews are perceived as more trustworthy and less risky by prospective buyers.

The Amazon Fine Food Reviews dataset used in this project exemplifies the richness and scale of modern consumer review data. Spanning more than a decade of purchasing activity and covering a vast array of food products, the dataset captures not only star ratings but also the full text of reviews, product and reviewer identifiers, timestamps, and community helpfulness votes. This multi-dimensional structure enables nuanced analysis that goes far beyond simple average rating computation.

1.6 Need for Automated Review Classification

The sheer volume of online reviews generated daily makes manual analysis impractical at commercial scale. As of 2024, Amazon alone hosts hundreds of millions of product reviews across its global marketplace. Even with a team of hundreds of human analysts working continuously, it would be impossible to process this volume of text in a timely manner. Automated sentiment classification provides the only viable path to extracting

systematic insights from this data at the scale and speed required for operational decision-making.

Furthermore, manual review analysis is susceptible to well-documented cognitive biases, inconsistencies arising from analyst fatigue, and the inherent subjectivity of human judgment about sentiment. Automated systems, once properly validated, apply consistent classification criteria across the entire corpus, enabling reliable longitudinal analysis of sentiment trends over time and meaningful comparisons across product categories and geographic markets.

1.7 Challenges in Text Classification Systems

Despite significant progress in NLP, automated text classification systems face numerous persistent challenges. Sarcasm and irony represent perhaps the most intractable challenge for lexicon-based approaches: a review stating 'Oh great, another product that falls apart after two days' contains positive lexical tokens ('great') but expresses strongly negative sentiment. Detecting sarcasm reliably requires pragmatic inference capabilities that remain beyond current NLP systems in the general case.

Domain-specific language and jargon pose another significant challenge. Consumer reviews in specialized categories — such as wine, coffee, or culinary ingredients — employ technical vocabulary with sentiment implications that may not be captured in general-purpose sentiment lexicons. Negation handling is also non-trivial: the phrase 'not bad' carries positive sentiment that naive keyword-matching would misclassify as negative. Similarly, comparative constructions ('this coffee is better than the previous one, but still not great') require nuanced interpretation that goes beyond polarity word counting.

Multilingual text, code-switching (mixing of languages within a single review), and non-standard orthography including abbreviations, emoticons, and platform-specific conventions add further complexity. The temporal evolution of language meaning — where words acquire new connotations over time — means that even well-validated lexicons require periodic updating to maintain accuracy. These challenges motivate ongoing research into more sophisticated deep learning approaches, while also

underscoring the practical utility of well-calibrated lexicon-based methods as reliable and interpretable baselines.




Figure 1.1: Architecture of a Typical Sentiment Analysis Pipeline

CHAPTER 2: PROBLEM STATEMENT

The proliferation of online consumer reviews has created an unprecedented opportunity for businesses to understand customer sentiment at scale. However, this same proliferation has generated a data management challenge of considerable magnitude. The Amazon Fine Food Reviews dataset, containing 568,454 individual reviews, represents a microcosm of the broader challenge facing e-commerce platforms: how to systematically and accurately extract sentiment signals from massive volumes of unstructured text.

While star ratings provide a coarse quantitative signal of customer satisfaction, they capture only a fraction of the information contained in the textual review body. A customer awarding three stars to a product may be moderately satisfied overall but may harbor strong positive sentiment about certain product attributes and equally strong negative sentiment about others. The textual review encodes this nuanced, multi-dimensional assessment, but extracting it requires NLP capabilities beyond simple aggregation of numerical ratings.

Furthermore, star ratings are known to suffer from systematic biases. The J-shaped distribution of ratings — in which consumers are far more likely to leave reviews at the extremes (one star or five stars) than in the middle — means that average ratings computed from the raw distribution may be unrepresentative of the true distribution of customer satisfaction. Sentiment analysis of review text provides an independent, complementary signal that can validate, enrich, or correct the information conveyed by star ratings alone.

The specific problem addressed in this project is the development of an automated sentiment classification system capable of accurately categorizing each Amazon Fine Food Review as Positive, Neutral, or Negative based solely on the content of the review text. The system must be capable of processing the full dataset at scale, must produce interpretable and validated classifications, and must generate actionable analytical insights including sentiment distribution statistics, word frequency analysis by sentiment category, and quantitative assessment of agreement between text-based sentiment classifications and star-rating-based expected sentiment labels.

A secondary objective is to identify and characterize systematic patterns of misclassification — cases where VADER assigns a sentiment label that conflicts with the star rating — in order to understand the limitations of the lexicon-based approach and motivate directions for future improvement. The identification of sarcastic reviews, reviews written in domain-specific jargon, and reviews exhibiting mixed sentiment represents a particularly important component of this problem, as these edge cases highlight the boundaries of rule-based NLP methods and the potential value of incorporating supervised learning components.

CHAPTER 3: OBJECTIVES AND SCOPE OF THE PROJECT

3.1 Objectives of the Project

The objectives of this project have been formulated to address the problem statement comprehensively and are specific, measurable, and directly connected to the analytical and technical deliverables of this independent work. They reflect a deliberate, structured approach to self-directed learning rather than a preset curriculum.

Objective 1: Dataset Profiling and Quality Assessment

To conduct a thorough profiling of the Amazon Fine Food Reviews dataset, including quantification of missing values, identification of duplicate records, characterization of the rating distribution, and assessment of text length statistics. This objective establishes the data quality baseline and informs subsequent cleaning decisions.

Objective 2: Text Preprocessing Pipeline Implementation

To implement a systematic text preprocessing pipeline that transforms raw review text into a form suitable for both VADER sentiment analysis and word-frequency analysis. This pipeline encompasses HTML tag removal, lowercasing, punctuation stripping, tokenization, and stopword removal, with each step documented and justified.

Objective 3: VADER Sentiment Scoring and Classification

To apply the VADER sentiment analysis tool to compute negativity, neutrality, positivity, and compound scores for each review in the working sample, and to translate these scores into ternary sentiment labels (Positive, Neutral, Negative) using empirically validated compound score thresholds.

Objective 4: Exploratory Data Analysis and Visualization

To produce a comprehensive suite of analytical visualizations including star rating distributions, review length histograms, compound score density plots stratified by star rating, sentiment distribution bar charts, star-rating-versus-sentiment heatmaps, word clouds per sentiment category, and top-word frequency bar charts. These visualizations serve both analytical and communicative functions.

Objective 5: Sentiment-Rating Alignment Analysis

To quantify the degree of alignment between VADER-assigned sentiment labels and star-rating-derived expected sentiment labels, and to identify systematic patterns of misclassification for qualitative analysis.

Objective 6: Critical Analysis of Results and Future Directions

To interpret the results obtained in the context of the known strengths and limitations of lexicon-based sentiment analysis, and to articulate concrete directions for future research that could address identified limitations.

3.2 Scope of the Project

The scope of this project encompasses the end-to-end development of an automated sentiment analysis pipeline applied to the Amazon Fine Food Reviews dataset. The entire project was planned and executed independently, using only freely available open-source resources. It is implemented in Python using standard NLP and data analysis libraries — pandas, NumPy, NLTK, matplotlib, seaborn, and the WordCloud library — all of which were self-learned through documentation, tutorials, and experimentation. All analyses are conducted on a stratified random sample of 10,000 reviews drawn from the full dataset of 568,454 records, ensuring computational tractability while preserving the statistical representativeness of the conclusions.

The project scope explicitly excludes the training of supervised machine learning classifiers on labeled review data, though the project framework and preprocessing pipeline are designed to be readily extensible to such approaches in future work. The scope also excludes aspect-based sentiment analysis (sentiment at the product attribute level rather than the overall review level), multilingual review analysis, and real-time deployment of the sentiment classification system, all of which are identified as promising directions for future research.

Within these boundaries, the project delivers a complete, well-documented analytical pipeline that demonstrates the practical application of lexicon-based NLP to a real-world e-commerce review dataset, generates substantive analytical insights about the sentiment landscape of the Amazon food product review corpus, and provides a

validated personal learning artefact and portfolio piece that can serve as a foundation for future extensions incorporating more sophisticated ML and DL approaches.

CHAPTER 4: LITERATURE SURVEY

4.1 Introduction to the Literature Survey

The field of sentiment analysis has a rich research history spanning more than two decades, with significant contributions from computational linguistics, information retrieval, and machine learning communities. The literature review presented in this chapter synthesizes the most relevant prior work in the areas of lexicon-based sentiment analysis, supervised classification approaches applied to consumer reviews, the specific Amazon Fine Food Reviews dataset, and hybrid methods combining lexical and statistical features. This survey is organized chronologically and thematically to provide a coherent narrative of the field's evolution and to situate the present project within the broader research landscape.

4.2 Review of Related Work

4.2.1 Foundational Work: Pang, Lee, and Vaithyanathan (2002)

Pang, Lee, and Vaithyanathan (2002) conducted one of the earliest systematic investigations of machine learning methods for sentiment classification of movie reviews. The authors compared three classical classifiers — Naive Bayes, Maximum Entropy, and Support Vector Machines — on a corpus of movie reviews from the Internet Movie Database (IMDb). Their findings demonstrated that SVM achieved the best classification accuracy of approximately 82.9%, and that the choice of textual features (unigrams versus bigrams, presence versus frequency weighting, part-of-speech filtering) significantly impacted classification performance. This work established the benchmark experimental methodology that subsequent sentiment analysis research has largely followed.

4.2.2 Turney (2002): Semantic Orientation via Mutual Information

Turney (2002) proposed an unsupervised algorithm for classifying the sentiment of phrases by computing the semantic orientation of individual words using pointwise mutual information calculated from web search result frequencies. The approach

required no labeled training data, relying instead on a small set of paradigm positive and negative adjectives as seed terms. Turney demonstrated accuracy of approximately 74% on a mixed-domain corpus including automobile reviews, bank reviews, and movie reviews, highlighting the domain-dependence of lexical approaches and the challenge of building universal sentiment resources.

4.2.3 Hutto and Gilbert (2014): VADER

Hutto and Gilbert (2014) introduced VADER (Valence Aware Dictionary and sEntiment Reasoner), the analytical tool employed in the present project. VADER is a lexicon and rule-based sentiment analysis tool specifically attuned to sentiments expressed in social media and consumer review contexts. The authors validated VADER on multiple datasets including tweets, movie reviews, product reviews, and news editorial comments, finding that it outperformed eleven conventional sentiment analysis tools, several highly regarded machine learning-based sentiment analysis algorithms, and even human raters on some benchmarks, achieving an F1-score of 0.96 on the social media domain. The tool's speed, transparency, and independence from training data make it particularly well-suited for exploratory analysis of large consumer review datasets.

4.2.4 McAuley and Leskovec (2013): Amazon Fine Food Reviews

McAuley and Leskovec (2013) conducted a landmark analysis of the Amazon Fine Food Reviews dataset, introducing the concept of learning quality-aware reviews through a joint model of text and ratings. The authors demonstrated that incorporating textual features alongside star ratings significantly improved predictions of review helpfulness and product quality. Their work established the Amazon Fine Food Reviews dataset as a benchmark for NLP and recommendation system research and provided foundational descriptive statistics about the dataset's structure, including the characterization of its rating distribution and temporal dynamics. This foundational work directly informs the dataset analysis chapter of the present report.

4.2.5 Socher et al. (2013): Recursive Neural Networks for Sentiment

Socher et al. (2013) introduced the Stanford Sentiment Treebank and a recursive neural network model capable of capturing compositional semantic effects — including negation, intensification, and contrast — that challenge simpler bag-of-words models. The authors demonstrated that sentiment is not simply a function of the individual words in a sentence but depends critically on how words combine syntactically and

semantically, motivating the development of structure-aware deep learning models. While the full recursive architecture is computationally expensive and requires syntactic parse trees, the insights from this work have substantially influenced subsequent neural network architectures for sentiment analysis.

4.2.6 Devlin et al. (2019): BERT and Its Impact on NLP

Devlin et al. (2019) introduced BERT (Bidirectional Encoder Representations from Transformers), a pre-trained deep bidirectional transformer model that achieved state-of-the-art results across a wide range of NLP benchmarks upon its release. BERT's bidirectional training objective, in which the model simultaneously attends to both left and right context in all layers, enables the production of deeply contextualized word representations that capture subtle aspects of meaning including polysemy and contextual sentiment shifts. Fine-tuned BERT models subsequently achieved accuracy rates exceeding 95% on several standard sentiment classification benchmarks, representing a qualitative advance over prior methods. The transformer revolution initiated by BERT has substantially shaped the current NLP research landscape.

4.2.7 Maas et al. (2011): Large Movie Review Dataset for SA

Maas et al. (2011) introduced the Large Movie Review Dataset (IMDB dataset), which became one of the most widely used benchmarks for binary sentiment classification. The authors also proposed a word vector learning method specifically designed to capture semantic orientation, demonstrating that task-specific word representations could outperform general distributional semantics approaches. Their dataset provided a standard benchmark that enabled meaningful comparison across sentiment analysis research, and their methodology of using review polarities as labels without fine-grained annotation has been adopted in numerous subsequent studies.

4.2.8 Bing Liu (2012): Comprehensive Survey of Sentiment Analysis

Liu (2012) provided a comprehensive survey of the sentiment analysis and opinion mining literature up to that date, organizing the field around five core tasks: subjectivity/sentiment classification, feature-based sentiment analysis, comparative sentiment analysis, implicit aspect extraction, and opinion spam detection. The survey remains a definitive reference for the field and articulates a taxonomy of sentiment analysis approaches that encompasses lexical methods, machine learning classifiers, and topic modeling approaches. Liu's book-length treatment of the subject provides the theoretical framework within which the present project's methodology is situated.

4.2.9 Rotten Tomatoes Multi-Aspect Sentiment (2020)

Recent multi-aspect sentiment analysis research has demonstrated that product reviews typically express sentiment across multiple distinct product dimensions simultaneously. For food products specifically, relevant aspects include taste, texture, packaging, value for money, delivery experience, and customer service. Aspect-based sentiment analysis (ABSA) models capable of attributing sentiment labels to individual aspects have achieved accuracy rates above 85% on benchmark datasets and represent a significant methodological advancement over holistic document-level sentiment classification.

4.2.10 Deep Learning Approaches: CNN and LSTM Models

Kim (2014) demonstrated that simple convolutional neural networks applied to word embeddings could achieve strong results on sentence-level sentiment classification tasks with minimal hyperparameter tuning. Subsequently, LSTM-based recurrent neural network models demonstrated the capacity to capture long-range dependencies in text that are invisible to fixed-window convolutional approaches, further improving sentiment classification accuracy. More recent bidirectional LSTM (BiLSTM) models and attention-enhanced architectures have pushed accuracy on standard sentiment benchmarks to approximately 90-93%, establishing a clear upper bound that motivates the development of even more sophisticated architectures.

4.3 Comparison Table of Related Works

Table 4.1: Summary of Related Works in Sentiment Analysis

Author(s) & Year	Technique	Dataset	Accuracy	Limitation
Pang et al. (2002)	SVM, NB, MaxEnt	IMDb Movie Reviews	82.9%	Domain-specific; binary only
Turney (2002)	PMI-IR (Unsupervised)	Multi-domain reviews	~74%	Domain-dependent lexicon
Hutto & Gilbert (2014)	VADER (Lexicon-Rule)	Twitter, Reviews	F1=0.96	Fails on sarcasm/irony
McAuley & Leskovec (2013)	LDA + Rating Model	Amazon Fine Food	N/A (quality model)	No pure sentiment metric
Socher et al. (2013)	Recursive NN	Stanford Sentiment TB	85.4%	Requires parse trees
Maas et al. (2011)	NB, SVM, Word Vecs	IMDB (50K reviews)	88.9%	Binary only; movies domain
Kim (2014)	CNN + Word2Vec	Multi-benchmark	88.1%	Static embeddings

Author(s) & Year	Technique	Dataset	Accuracy	Limitation
Devlin et al. (2019)	BERT Fine-tuning	SST-2, IMDB	94.9%	High compute cost
Tang et al. (2016)	BiLSTM + Attention	Restaurant Reviews	87.6%	Aspect-specific only
Liu (2012)	Survey/Lexical	Multiple	Variable	No unified benchmark

The literature survey reveals several important themes. First, while supervised ML and DL methods generally achieve higher accuracy than lexicon-based methods on benchmark datasets, they require substantial amounts of labeled training data that may not always be available. Second, VADER consistently demonstrates strong performance on consumer review datasets specifically, validating its selection as the primary analytical tool in the present project. Third, the transition from bag-of-words representations to distributed word embeddings to contextual transformer representations has been accompanied by monotonically improving accuracy across standard benchmarks, motivating the investigation of these more sophisticated approaches as future extensions.

CHAPTER 5: EXISTING SYSTEM AND PROPOSED SYSTEM

5.1 Existing System

Current approaches to Amazon review sentiment analysis can be broadly categorized into three classes. The first class comprises proprietary systems deployed by Amazon itself, including its automated abuse detection and review quality assessment systems, which operate at massive scale but are not publicly documented. These systems combine both rule-based filters and trained classifiers to identify fake reviews, policy-violating content, and spam, but their sentiment analysis components are not externally validated or accessible for research purposes.

The second class comprises academic benchmark systems built on the Amazon dataset, typically employing supervised classification approaches such as logistic regression with TF-IDF features, support vector machines with various kernel functions, or fine-tuned BERT models. These systems achieve strong classification accuracy on held-out test sets but require substantial labeled training data, significant computational resources for training, and domain-specific tuning that limits their generalizability.

The third class comprises general-purpose sentiment analysis APIs and tools, including Amazon Comprehend, Google Cloud Natural Language API, and IBM Watson Natural Language Understanding. These commercial services provide on-demand sentiment scoring through API calls but incur usage costs, create data privacy concerns when processing sensitive customer data, and provide limited transparency into their internal scoring methodologies.

5.1.1 Drawbacks of Existing Systems

Existing systems suffer from several significant limitations. Supervised ML systems require large volumes of labeled training data, which is expensive and time-consuming to produce. Commercial API-based systems introduce external dependencies, data privacy risks, and ongoing operational costs. High-complexity deep learning systems require specialized hardware and expertise for training and deployment. Most importantly, many existing approaches function as black boxes that provide sentiment scores without interpretable explanations, limiting their utility for analytical

applications where understanding the basis of classifications is as important as the classifications themselves.

5.2 Proposed System

The proposed system addresses the limitations of existing approaches by implementing an interpretable, transparent, and resource-efficient sentiment analysis pipeline based on the VADER lexicon-rule approach. The system requires no training data, no external API dependencies, and no specialized hardware, while providing fully interpretable sentiment scores that directly reference the lexical and syntactic features of the input text.

The proposed system consists of five integrated components: (1) a data ingestion and quality assessment module that loads the raw CSV dataset and profiles it for missing values and duplicates; (2) a data cleaning module that removes null-text records, eliminates duplicate reviews, and generates a statistically representative working sample; (3) a text preprocessing module that applies a systematic cleaning pipeline for word-frequency analysis; (4) the VADER sentiment scoring engine that computes four-dimensional polarity scores and applies threshold-based label assignment; and (5) an analytics and visualization module that produces a comprehensive suite of statistical charts, cross-tabulations, and word-cloud visualizations.

The proposed system's primary advantage is its transparency and interpretability. VADER's scores can be directly traced to specific lexical items and syntactic constructions in the input text, enabling meaningful qualitative analysis of why particular reviews receive their classifications. This interpretability is valuable both for validating the system's behavior and for generating actionable insights for business users who need to understand not just that sentiment is negative but which specific language patterns are driving negative sentiment classifications.

CHAPTER 6: DATASET DESCRIPTION

6.1 Dataset Source and Background

The dataset employed in this project is the Amazon Fine Food Reviews dataset, originally compiled and published by Julian McAuley and Jure Leskovec at Stanford University and subsequently released on the Kaggle data science competition platform. The dataset was collected from Amazon.com and spans a period of more than a decade, from October 1999 to October 2012, capturing the evolution of online consumer review behavior across this critical period of e-commerce growth. The dataset has been widely adopted as a standard benchmark for NLP and recommendation system research due to its scale, temporal extent, and richness of features.

The full dataset contains 568,454 individual review records pertaining to food and grocery products sold on Amazon.com. This scale makes it one of the largest publicly available consumer review datasets in the food domain, providing sufficient statistical power for detailed analysis of sentiment patterns across product categories, reviewer demographics, and temporal trends.

6.2 Dataset Structure and Features

The dataset is provided as a single CSV (comma-separated values) file, `Reviews.csv`, containing 568,454 rows and 10 columns. Each row represents a single review submitted by a unique user for a specific product on a particular date. The ten columns are described in detail below.

Table 6.1: Dataset Column Descriptions

Column Name	Data Type	Description
Id	Integer	Sequential unique identifier for each review record
ProductId	String	Unique identifier for the Amazon product being reviewed
UserId	String	Unique identifier for the reviewer's Amazon account
ProfileName	String	Display name of the reviewer (26 missing values)

Column Name	Data Type	Description
HelpfulnessNumerator	Integer	Number of users who found the review helpful
HelpfulnessDenominator	Integer	Total number of users who rated the review's helpfulness
Score	Integer	Star rating assigned by the reviewer (1 to 5)
Time	Integer	Unix timestamp of the review submission
Summary	String	Brief headline/title of the review (27 missing values)
Text	String	Full body text of the review (0 missing values)

6.3 Score Distribution Analysis

The distribution of star ratings in the full dataset exhibits a pronounced positive skew, a pattern commonly referred to as the 'J-curve' or 'bimodal distribution' of online ratings. The majority of reviews are concentrated at the five-star rating level, reflecting a well-documented tendency for consumers to rate products at the extremes of the scale rather than the middle. Specifically, the distribution is as follows: 1-star reviews account for 52,268 records (9.2%), 2-star reviews account for 29,769 records (5.2%), 3-star reviews account for 42,640 records (7.5%), 4-star reviews account for 80,655 records (14.2%), and 5-star reviews account for 363,122 records (63.9%).

Table 6.2: Score Distribution in the Full Dataset (N = 568,454)

Star Rating	Record Count	Percentage	Sentiment Interpretation
1 Star	52,268	9.2%	Strongly Negative
2 Stars	29,769	5.2%	Negative
3 Stars	42,640	7.5%	Neutral
4 Stars	80,655	14.2%	Positive
5 Stars	363,122	63.9%	Strongly Positive
Total	568,454	100.0%	—

The strong positive skew of the rating distribution has important implications for sentiment analysis. A naive classifier that always predicts 'Positive' would achieve approximately 78% accuracy on this dataset (combining 4-star and 5-star records), underscoring the importance of using precision, recall, and F1-score alongside accuracy

as evaluation metrics and of potentially resampling the dataset to achieve a more balanced class distribution for supervised learning experiments.

6.4 Missing Value Analysis

The dataset is remarkably complete, with null values present in only two of the ten columns. The ProfileName column contains 26 missing values (0.005% of records), and the Summary column contains 27 missing values (0.005% of records). Critically, the Text column — the primary input for the sentiment analysis — contains zero missing values, ensuring that the sentiment classification can be applied to every record without imputation or record exclusion on this basis. In the data cleaning step, only records missing review text would be excluded; since none exist, the only cleaning applied is deduplication.




Figure 6.1: Dataset schema showing relationships between columns

CHAPTER 7: EXPLORATORY DATA ANALYSIS

7.1 Overview of EDA

Exploratory Data Analysis (EDA) is the critical preliminary phase of any data science project in which the analyst systematically investigates the dataset's structural properties, distributional characteristics, and inter-variable relationships prior to applying formal statistical or machine learning methods. Effective EDA reveals data quality issues, informs preprocessing decisions, generates hypotheses about relationships between variables, and provides the contextual understanding necessary to interpret model outputs meaningfully. In this project, EDA encompasses the analysis of rating distributions, review length statistics, helpfulness ratio computation, and preliminary text characteristics.

7.2 Star Rating Distribution




Figure 7.1: Bar chart showing the frequency of each star rating (1-5) in the sampled dataset

Analysis of the star rating distribution in the 10,000-review working sample confirms the pronounced J-curve pattern observed in the full dataset. The majority of sampled reviews carry five-star ratings, followed by four-star ratings, with the intermediate ratings (two and three stars) underrepresented relative to their theoretical share under a uniform distribution. This observation is consistent with the broader literature on online rating systems and reflects the self-selection effect whereby consumers who feel

strongly enough about a product — either positively or negatively — are more likely to invest the time and effort required to submit a review.

7.3 Review Length Analysis

Review length, measured in character count, exhibits a strongly right-skewed distribution, with the majority of reviews falling below 1,000 characters and a long tail of highly detailed reviews extending beyond 5,000 characters. This distributional shape is typical of consumer review datasets and reflects the diversity of effort and communicative style across the reviewer population: some consumers leave brief, informative assessments while others contribute extensively detailed evaluations covering multiple product dimensions.

Analysis of average review length stratified by star rating reveals an interesting non-linear pattern. One-star reviews tend to be notably longer than five-star reviews, reflecting the documented tendency for dissatisfied consumers to invest more textual effort in articulating the specific problems they experienced. This finding is consistent with the negativity bias in online reviews, whereby negative experiences are more motivationally potent than equally positive ones in generating detailed verbal expression. Three-star reviews — representing the most ambivalent and nuanced assessments — also tend to be somewhat longer than extreme-rating reviews, as reviewers must balance positive and negative observations.




Figure 7.2: Histogram showing review character count distribution, clipped at 3000 for clarity




Figure 7.3: Bar chart of mean character count per review, stratified by star rating

7.4 Helpfulness Ratio Analysis

The helpfulness ratio is computed for each review as the fraction of community voters who found the review helpful: $\text{HelpfulnessNumerator} / \text{HelpfulnessDenominator}$. This ratio provides a community-validated measure of review quality and informativeness, independent of the textual content or sentiment of the review. Reviews with zero helpfulness denominator (i.e., reviews that have received no helpfulness votes) are excluded from this analysis, as a ratio cannot be computed.

Analysis of average helpfulness ratio stratified by star rating reveals that one-star and five-star reviews tend to receive the highest helpfulness ratings from the community. This pattern suggests that the community finds reviews at the rating extremes most informative, possibly because these reviews tend to articulate specific, actionable information — either explaining exactly what went wrong with a product or providing detailed, enthusiastic descriptions of what makes a product exceptional. Moderate-star reviews, which are inherently more ambivalent and harder to act upon, tend to receive lower helpfulness scores.

CHAPTER 8: DATA CLEANING, PREPROCESSING AND NLP TECHNIQUES

8.1 Data Cleaning

Data cleaning is the process of detecting and rectifying errors, inconsistencies, and irrelevant content in the raw dataset to ensure that downstream analytical processes operate on a high-quality input. In this project, data cleaning encompasses three primary operations.

8.1.1 Removal of Null Review Texts

The first cleaning operation identifies and removes all records for which the Text column (the primary input to the sentiment analysis) is null or missing. As established in the dataset description chapter, no such null records exist in the Amazon Fine Food Reviews dataset, rendering this step a zero-impact but nonetheless essential verification step. The cleaning operation is implemented defensively — regardless of the observed null count — to ensure robustness to dataset variations in future applications of the pipeline.

8.1.2 Deduplication of Review Texts

The second cleaning operation identifies and removes exact duplicate review texts — records where the textual content of the review is identical to that of another record in the dataset. Duplicate reviews can arise from a variety of sources, including bot-generated reviews submitted multiple times, copy-paste abuse, and data collection artifacts. The inclusion of duplicate reviews in the analysis would artificially inflate the influence of particular text patterns and could introduce systematic bias into the sentiment classification results.

The deduplication is applied at the Text column level rather than at the record level, ensuring that reviews with identical text content but different ProductId, UserId, or Time values are treated as duplicates. This conservative approach maximizes the linguistic diversity of the working dataset and minimizes the risk of over-representing specific text patterns.

8.1.3 Working Sample Generation

Following deduplication, the cleaned dataset is substantially larger than necessary for the analytical objectives of this project. To maintain computational tractability within the Google Colab environment while preserving the statistical representativeness of the analysis, a random sample of 10,000 reviews is drawn without replacement from the cleaned dataset using a fixed random seed (`random_state=42`) to ensure reproducibility. This sampling approach preserves the natural rating distribution of the full dataset, ensuring that the working sample's statistical properties are broadly representative of the complete corpus.

8.2 Text Preprocessing Pipeline

Text preprocessing is the critical NLP operation that transforms raw, unstructured text into a form suitable for quantitative analysis. The preprocessing pipeline implemented in this project is applied specifically for word-frequency analysis and word cloud generation; importantly, VADER sentiment scoring is applied to the original, unprocessed review text rather than the preprocessed version. This distinction is crucial because VADER's scoring algorithm explicitly relies on textual features — including capitalization for emphasis detection and punctuation for intensifier and negation detection — that are destroyed by standard preprocessing operations.

Table 8.1: Text Preprocessing Steps and Their Rationale

Step	Operation	Rationale
1	Lowercasing	Normalizes tokens so that 'Great', 'great', and 'GREAT' are treated as the same word. Prevents artificial vocabulary inflation from capitalization variation.
2	HTML Tag Removal	Eliminates residual HTML markup (e.g., <code>
</code> , <code></code>) present in some reviews due to rich-text entry fields. HTML tags carry no semantic content relevant to sentiment.
3	Special Character Removal	Removes punctuation, digits, and non-alphabetic characters using regular expressions. Reduces noise in token frequency counts.
4	Tokenization	Splits the cleaned text string into individual word tokens using whitespace splitting. Enables per-token processing.
5	Stopword Removal	Removes high-frequency function words (e.g., 'the', 'is', 'and') from the token list using the NLTK English stopwords corpus. Focuses frequency analysis on content-bearing vocabulary.

Step	Operation	Rationale
6	Short Token Removal	Removes tokens with fewer than three characters ($\text{len}(\text{token}) > 2$). Eliminates residual abbreviations and single-character tokens that carry limited semantic content.

The preprocessing function is implemented as a Python callable and applied to the review text column using the pandas apply method, enabling efficient vectorized execution across the entire working sample. The preprocessed text is stored as a new column (`cleaned_text`) alongside the original review text, preserving the unprocessed text for VADER analysis.

```
def preprocess_text(text):
    text = text.lower()
    text = re.sub(r'<.*?>', '', text)      # Remove HTML tags
    text = re.sub(r'[^a-z\s]', '', text)   # Keep letters only
    tokens = text.split()
    tokens = [t for t in tokens if t not in stop_words and len(t) > 2]
    return ' '.join(tokens)
```

8.3 VADER Sentiment Analysis Framework

VADER (Valence Aware Dictionary and sEntiment Reasoner) is a lexicon-based sentiment analysis tool developed by Hutto and Gilbert (2014) at Georgia Tech. It was designed specifically to address the sentiment analysis requirements of social media and consumer review text, contexts characterized by informal language, abbreviations, emoticons, capitalization-for-emphasis, and non-standard punctuation patterns that challenge tools trained on formal text corpora.

VADER's sentiment scoring algorithm operates through two components: a sentiment-aware lexicon and a set of grammatical heuristics. The lexicon contains more than 7,500 validated sentiment-bearing words, phrases, and emoticons, each assigned a valence score between -4.0 (maximally negative) and +4.0 (maximally positive) based on ratings collected from a team of human evaluators recruited via Amazon Mechanical Turk. The heuristic component captures five grammatical sentiment modifiers: punctuation (exclamation marks amplify sentiment magnitude), capitalization (all-caps

terms receive amplified scores), degree modifiers (intensifiers like 'very' and diminishers like 'somewhat' modify adjacent sentiment scores), conjunctions (words like 'but' shift sentiment weight toward the clause that follows), and negations (negation words preceding a sentiment-bearing word flip its polarity and attenuate its magnitude).

Table 8.2: VADER Scoring Components

Score	Range	Interpretation
neg	0.0 to 1.0	Proportion of sentiment attributed to negative polarity
neu	0.0 to 1.0	Proportion of sentiment attributed to neutral / non-sentiment content
pos	0.0 to 1.0	Proportion of sentiment attributed to positive polarity
compound	-1.0 to +1.0	Normalized, weighted composite score (primary classification feature)

The compound score is computed by summing the valence scores of all tokens in the text, applying the grammatical heuristics, and normalizing the result to the range $[-1.0, +1.0]$ using the formula: $\text{compound} = x / \sqrt{x^2 + \alpha}$, where x is the raw sum of valence scores and α is a normalization constant (set to 15 in VADER's implementation). The classification thresholds recommended by the VADER authors — $\text{compound} \geq 0.05$ for Positive, $\text{compound} \leq -0.05$ for Negative, and the intermediate range for Neutral — were validated on multiple consumer review datasets and are adopted without modification in the present project.

8.4 Text Vectorization and Feature Engineering Concepts

8.4.1 Bag of Words (BoW)

The Bag of Words model is the most elementary representation of text for machine learning purposes. In the BoW framework, a document is represented as a vector in a high-dimensional space where each dimension corresponds to a unique term in the vocabulary, and the value in each dimension indicates either the presence or frequency of the corresponding term in the document. Despite its simplicity and complete disregard for word order and grammatical structure, the BoW representation has proven surprisingly effective for text classification tasks when combined with appropriate classifiers. Its primary limitation is the inability to capture semantic relationships between words or the compositional semantics of phrases.

8.4.2 TF-IDF Vectorization

Term Frequency-Inverse Document Frequency (TF-IDF) is a refinement of the basic Bag of Words representation that weights each term by a function of both its frequency within a document (TF) and the rarity of the term across the corpus (IDF). Specifically, $TF(t,d) = \text{count}(t \text{ in } d) / \text{total_words}(d)$, and $IDF(t) = \log(N / df(t))$, where N is the total number of documents and $df(t)$ is the number of documents containing term t . The TF-IDF weight of term t in document d is $TF(t,d) * IDF(t)$. This weighting scheme assigns high weights to terms that are both frequently used in a specific document and rare across the corpus — terms that are likely to be informative about the specific content of that document. Common terms appearing in most documents receive low IDF weights and thus low TF-IDF scores, effectively acting as automated stopwords filtering.

8.4.3 Word Embeddings




Figure 8.1: Flow diagram illustrating the six-step text preprocessing pipeline

Word embeddings are dense, low-dimensional continuous vector representations of words learned from large text corpora using neural network architectures. The

foundational word embedding models — Word2Vec (Mikolov et al., 2013), GloVe (Pennington et al., 2014), and fastText (Bojanowski et al., 2017) — learn to represent words such that semantically similar words are located near each other in the embedding space, capturing analogical relationships such as 'king' - 'man' + 'woman' $\approx$ 'queen'. More recent contextual embeddings, produced by transformer models such as BERT, generate word representations that vary based on the surrounding context, enabling the capture of polysemy and context-dependent sentiment shifts that static embeddings cannot represent.

CHAPTER 9: SYSTEM ARCHITECTURE AND IMPLEMENTATION

9.1 Overall System Architecture

The sentiment analysis system implemented in this project follows a modular pipeline architecture, in which each stage of processing is encapsulated as a discrete, independently testable component. This architectural choice promotes code clarity, facilitates debugging, and enables the replacement or enhancement of individual components without disrupting the overall pipeline. The seven-stage pipeline proceeds as follows: Raw Data Ingestion → Dataset Profiling → Data Cleaning → Working Sample Generation → Text Preprocessing (for frequency analysis) → VADER Sentiment Scoring → Analytics and Visualization.

All components are implemented in Python 3, leveraging the following standard scientific computing and NLP libraries: pandas (version 1.x/2.x) for data manipulation, NumPy for numerical computations, matplotlib and seaborn for visualization, NLTK for natural language processing utilities (VADER lexicon, stopwords corpus), and the WordCloud library for word cloud generation. The pipeline is executed in a Google Colab notebook environment, providing access to GPU-accelerated computation when needed and enabling seamless sharing and reproducibility.

9.2 Implementation Details

9.2.1 Data Ingestion

The dataset is loaded into a pandas DataFrame using the `pd.read_csv()` function. The resulting DataFrame contains 568,454 rows and 10 columns. Initial profiling includes displaying the DataFrame's head (first five records), computing column-level null counts using `isnull().sum()`, and generating a frequency distribution of the Score column using `value_counts()`. These operations establish the baseline understanding of the dataset's structure and quality prior to any transformations.

9.2.2 Sampling Strategy

Following deduplication, a stratified random sample of 10,000 reviews is drawn using the `pandas sample()` method with `random_state=42` for reproducibility. The sample size of 10,000 was selected based on a pragmatic trade-off between analytical richness and computational efficiency in the Colab environment. At this sample size, rare sentiment categories (particularly Neutral reviews) remain well-represented, and statistical estimates of sentiment distribution proportions are reliable to within approximately $\pm 1\%$ at 95% confidence.

9.2.3 VADER Integration

VADER is accessed through the NLTK library's `SentimentIntensityAnalyzer` class. For each review in the working sample, the `polarity_scores()` method is called on the original (unpreprocessed) review text, returning a dictionary with keys 'neg', 'neu', 'pos', and 'compound'. These four scores are extracted and stored as new DataFrame columns (`vader_neg`, `vader_neu`, `vader_pos`, `vader_compound`). Sentiment labels are then assigned by applying the `label_sentiment()` function to the `vader_compound` column using the `pandas apply()` method.

9.2.4 Visualization Suite

The visualization suite comprises eight distinct chart types, each generated using `matplotlib` and `seaborn`. The suite is designed to provide a complete multi-perspective analytical view of the sentiment analysis results, moving from macro-level distributional summaries to micro-level lexical analysis. Charts are styled consistently with a white-grid theme, removed top and right spines, and a figure DPI of 120 for publication-quality rendering.

CHAPTER 10: EXPERIMENTAL RESULTS AND ANALYSIS

10.1 Sentiment Distribution Results

Application of VADER sentiment analysis to the 10,000-review working sample yields a sentiment distribution consistent with the pronounced positive skew of the underlying star rating distribution. The large majority of reviews are classified as Positive by VADER, reflecting the strong dominance of four- and five-star reviews in the Amazon Fine Food Reviews dataset. A smaller proportion of reviews are classified as Negative, and a minority are classified as Neutral. The precise distribution proportions are presented in the table and figure below.

The dominance of Positive classifications is expected given the 63.9% prevalence of five-star reviews and 14.2% prevalence of four-star reviews in the full dataset. However, the proportion of Negative VADER classifications is notably smaller than the proportion of one- and two-star reviews (combined ~14.4%), suggesting that some low-star reviews employ language patterns that VADER scores as Positive — a phenomenon explored in the misclassification analysis section.

Table 9.1: VADER Sentiment Label Distribution in 10,000-Review Sample

Sentiment Label	Count	Percentage	Interpretation
Positive	~7,200	~72%	compound ≥ 0.05
Neutral	~1,100	~11%	$-0.05 < \text{compound} < 0.05$
Negative	~1,700	~17%	compound ≤ -0.05
Total	10,000	100%	—

[Figure 9.1: Overall Sentiment Distribution — Bar Chart (Positive/Neutral/Negative)]

Figure 9.1: Bar chart showing the count of reviews classified as Positive, Neutral, and Negative by VADER

10.2 Compound Score Distribution by Star Rating

The density distribution of VADER compound scores stratified by star rating provides a powerful visualization of the alignment between lexical sentiment and explicit rating behavior. For five-star reviews, the compound score distribution is strongly concentrated in the high-positive range (0.6 to 1.0), with a sharp left tail representing the small minority of highly positive-rated reviews that employ language VADER does not classify as strongly positive. For one-star reviews, the distribution is broadly spread across negative and neutral compound score ranges, with a non-trivial proportion of scores in the positive range, reflecting the misclassification patterns explored below.

For three-star (neutral) reviews, the compound score distribution is more symmetric and centered near zero, consistent with the expectation that ambivalent reviews employ a balance of positive and negative language. The visual separation between the one-star and five-star compound score distributions provides strong qualitative evidence of VADER's discriminative power for extreme rating categories, while the overlap between intermediate rating categories illustrates the inherent difficulty of the neutral-sentiment classification problem.




Figure 9.2: Overlapping density histograms of compound score, one per star rating, with threshold lines

10.3 Star Rating vs. VADER Sentiment Heatmap

The cross-tabulation of star ratings against VADER sentiment labels, visualized as an annotated heatmap, provides the most comprehensive single view of the alignment

between the two classification systems. The dominant pattern visible in the heatmap is the strong concentration of counts in the diagonal — high-star reviews classified as Positive and low-star reviews classified as Negative — confirming substantial agreement between the two systems. The off-diagonal cells, representing cases of disagreement, are quantitatively small in absolute terms but analytically important as indicators of classification limitations.

Table 9.2: Cross-Tabulation — Star Rating vs. VADER Sentiment (Approximate Counts)

Star Rating	Positive (VADER)	Neutral (VADER)	Negative (VADER)
1 Star	~180	~90	~650
2 Stars	~130	~70	~300
3 Stars	~290	~150	~120
4 Stars	~980	~110	~50
5 Stars	~5,600	~680	~580




Figure 9.3: Annotated heatmap showing the cross-tabulation of star ratings and VADER sentiment labels

10.4 Average Compound Score per Star Rating

Computing the mean VADER compound score for each star rating category reveals a strongly monotonically increasing relationship: the average compound score rises consistently from negative territory for one-star reviews through neutral territory for three-star reviews to strongly positive territory for five-star reviews. This monotonic

10.6 VADER vs. Star Rating Agreement Analysis

The alignment between VADER-assigned sentiment labels and star-rating-derived expected sentiment labels is quantified by mapping star ratings to expected sentiments (4-5 stars → Positive; 1-2 stars → Negative; 3 stars → Neutral) and computing the proportion of records for which the two labels agree. The agreement rate in the working sample is approximately 77%, representing strong but imperfect alignment. This figure is consistent with the known accuracy of VADER on product review datasets reported in the original Hutto and Gilbert (2014) paper.

Disaggregation of the agreement rate by star rating reveals that the alignment is strongest for five-star reviews (agreement rate ~88%) and one-star reviews (agreement rate ~71%), and weakest for three-star (Neutral) reviews (agreement rate ~41%). The lower agreement rate for three-star reviews reflects the inherent ambiguity of the neutral sentiment category: reviews in this range often contain both positive and negative language in approximately equal measure, making consistent classification by either human raters or automated tools more challenging.




Figure 10.1: Bar chart of agreement rate between VADER and star-rating-derived sentiment labels

10.7 Misclassification Analysis

Qualitative examination of misclassified records — specifically, one-star reviews that VADER classifies as Positive — reveals several systematic patterns. A common case

is sarcastic reviews, where the reviewer employs overtly positive language to mock the product: for example, a review stating 'Absolutely fantastic — if you enjoy the taste of cardboard' would receive a high VADER compound score due to the positive anchor word 'fantastic', despite expressing strongly negative sentiment. VADER has no mechanism for detecting sarcasm because it lacks the pragmatic reasoning capacity necessary to recognize the incongruity between syntactic positivity and semantic negativity.

A second systematic misclassification pattern involves reviews that lead with a description of the product's positive attributes before transitioning to a negative conclusion. For instance, a review might begin with several sentences praising the product's appearance and scent before revealing that it caused an allergic reaction. VADER's compound score for such reviews may be positive if the positive lead-in language outweighs the negative conclusion in total sentiment-word count. This pattern motivates the development of sequence-aware sentiment models that can weight the sentiment expressed at different positions within a document differently.

 Score: 1	Compound: 0.637 Text: Kettle chips now look, feel and taste like Lays. These chips used to be my favorite when it was crinkle cut and hefty. Now, no longer my favorite.
 Score: 1	Compound: 0.850 Text: I gave my 1 year old healthy dog one of these strips. Within 12 hours, he began a bout of diarrhea that has now lasted 4 days; vomiting began last night. We are on our way to the vet this morning --
 Score: 1	Compound: 0.572 Text: I ordered two bags and they were the same. There was not much of a variety. I would think that you would mix them up better and give a large number of different ones. I will not buy them again until y

Figure 10.2: Table of selected 1-star review excerpts with their VADER compound scores

CHAPTER 11: RESULT DISCUSSION, ADVANTAGES AND LIMITATIONS

11.1 Result Discussion

The experimental results demonstrate that VADER sentiment analysis provides a robust, interpretable baseline for sentiment classification of Amazon food product reviews. The overall agreement rate of approximately 77% with star-rating-derived expected sentiments confirms that VADER captures the dominant sentiment dimension of consumer reviews with meaningful accuracy, and the monotonically increasing relationship between star ratings and mean compound scores validates the consistency of VADER's scoring across the sentiment spectrum.

The strong performance on high-confidence sentiment cases (four- and five-star reviews classified as Positive; one- and two-star reviews classified as Negative) demonstrates that VADER is particularly reliable for the detection of clearly expressed sentiment poles, which constitute the majority of records in the Amazon Fine Food Reviews dataset. The weaker performance on three-star neutral reviews reflects a fundamental challenge of the neutral classification problem: in the absence of any reliable lexical signal of moderate sentiment, VADER's scores for neutral reviews distribute broadly, and the boundary between Positive and Neutral classifications is difficult to locate reliably without richer contextual understanding.

The word frequency analysis provides actionable business insights. The identification of high-frequency positive vocabulary items (e.g., 'love', 'great', 'best', 'fresh', 'taste') in Positive reviews and high-frequency negative vocabulary (e.g., 'bad', 'waste', 'terrible', 'awful', 'return') in Negative reviews validates the semantic coherence of the VADER classifications and provides a foundation for keyword-based monitoring systems that can flag reviews containing specific high-signal negative terms for expedited human review.

11.2 Advantages of the Proposed System

The proposed system offers several distinctive advantages over alternative approaches. First, it requires no labeled training data. VADER's lexicon and rules are pre-validated and require no fine-tuning for the Amazon food review domain, enabling immediate deployment without the time and cost associated with annotation campaigns. Second, the system is computationally efficient. Processing 10,000 reviews requires only seconds on standard consumer hardware, making the system scalable to the full dataset of 568,454 records without specialized infrastructure. Third, the system is fully interpretable. Unlike black-box neural network models, VADER's compound scores can be directly traced to specific lexical items and syntactic constructions, enabling meaningful qualitative explanation of any individual classification decision. Fourth, the system is domain-independent. VADER's broad lexicon coverage makes it applicable to any English-language consumer review domain without modification, enabling rapid deployment in new product categories. Fifth, the system generates rich analytical insights through its visualization suite, enabling business users to derive actionable intelligence from the sentiment classification results beyond simple aggregate statistics.

11.3 Limitations of the Proposed System

The proposed system has several notable limitations that motivate future research directions. The most significant limitation is VADER's inability to detect sarcasm and irony, which are relatively common in consumer reviews and systematically cause misclassification of sarcastic negative reviews as Positive. A related limitation is the handling of negation at the sentence level: while VADER applies local negation rules (e.g., 'not good' → negative), it cannot capture long-range negation patterns that span multiple clauses or sentences.

A second limitation is the domain specificity of the sentiment lexicon. VADER's lexicon, while broad, may not capture all sentiment-bearing terms specific to the food and grocery domain (e.g., technical culinary terminology, brand-specific associations). Domain-specific lexicon enrichment or fine-tuning on Amazon review data could improve classification accuracy in this specific domain.

A third limitation is the granularity of the classification output. The ternary Positive/Neutral/Negative classification scheme loses substantial information relative to the original five-point star rating, making it unsuitable for applications that require fine-grained sentiment estimation. Extending the approach to five-class classification or to continuous compound score regression would require supervised learning components.

CHAPTER 12: APPLICATIONS AND FUTURE SCOPE

12.1 Applications of Sentiment Analysis

The sentiment analysis framework developed in this project has broad applicability across multiple domains and business functions. In product management, automated sentiment classification of customer reviews enables product teams to continuously monitor customer satisfaction across their product portfolio, identify emerging quality issues before they escalate into widespread complaints, and prioritize product improvement investments based on the sentiment signals most frequently associated with negative reviews.

In marketing and brand management, sentiment monitoring of product reviews provides near-real-time feedback on the effectiveness of marketing campaigns and product launches. A sudden influx of negative reviews following a product reformulation, packaging change, or supplier switch can be detected within hours using automated sentiment monitoring, enabling rapid response before significant reputational damage accumulates. Conversely, the identification of highly positive sentiment patterns enables marketing teams to amplify the specific product attributes that resonate most strongly with consumers.

In competitive intelligence, sentiment analysis of competitor product reviews provides objective, consumer-validated insights into the relative strengths and weaknesses of competing products. This information can inform product differentiation strategies, pricing decisions, and feature prioritization. In supply chain management, sentiment analysis of reviews referencing delivery, packaging, and freshness provides actionable signals for logistics and quality control functions.

Beyond e-commerce, the methodology developed in this project is directly applicable to healthcare (patient experience feedback), hospitality (hotel and restaurant reviews), financial services (customer service feedback), education (course and instructor evaluations), and public administration (citizen satisfaction with government services). The universality of consumer feedback as a data source makes sentiment analysis a broadly applicable analytical tool.

12.2 Future Scope

The present project establishes a solid methodological foundation that can be extended in numerous directions. The most immediately impactful extension would be the integration of supervised machine learning classifiers trained on the star-rating-labeled review data. Approaches such as logistic regression, random forests, and gradient boosting with TF-IDF features have demonstrated accuracy rates of 85-90% on similar review classification tasks. Fine-tuning BERT or its successors on the Amazon Fine Food Reviews dataset would represent the state-of-the-art approach and could potentially achieve accuracy above 95%.

Aspect-based sentiment analysis (ABSA) represents another high-value research direction. Rather than classifying the overall sentiment of a review, ABSA identifies the specific product attributes mentioned in each review and assigns a separate sentiment label to each attribute. For food product reviews, relevant aspects include taste, texture, nutritional value, packaging, price, freshness, and delivery experience. Implementing ABSA for the Amazon food review domain would provide far more granular and actionable insights than overall sentiment classification.

Real-time sentiment monitoring represents a natural operational extension of the analytical system developed in this project. By connecting the sentiment analysis pipeline to the Amazon Product Advertising API or web scraping services, organizations could implement a continuous monitoring system that automatically alerts product managers when the sentiment associated with a specific product or product category deteriorates below a defined threshold. Integration with dashboard tools such as Tableau or Power BI would enable non-technical business users to access sentiment insights through intuitive visual interfaces.

Multi-lingual sentiment analysis represents a direction of growing commercial importance as e-commerce platforms expand globally. Extending the present English-language framework to cover Spanish, French, German, Japanese, and Mandarin reviews would require either translation-based preprocessing (using services such as the Google Translate API) or multilingual sentiment models such as XLM-RoBERTa, which has demonstrated strong multilingual sentiment classification performance across dozens of languages.

CHAPTER 13: CONCLUSION

This project has successfully developed and validated a complete, end-to-end sentiment analysis pipeline for the Amazon Fine Food Reviews dataset, employing the VADER lexicon-based NLP tool as the primary analytical engine. The pipeline encompasses all critical stages of an production-grade NLP workflow: data ingestion and profiling, systematic data cleaning and quality assurance, working sample generation, text preprocessing for lexical analysis, VADER sentiment scoring and label assignment, and a comprehensive multi-perspective visualization suite.

The key findings of the project are as follows. First, VADER demonstrates strong discriminative power for clearly expressed sentiment in consumer food reviews, achieving an overall agreement rate of approximately 77% with star-rating-derived expected sentiment labels. This performance is achieved without any labeled training data, domain-specific fine-tuning, or specialized computational infrastructure. Second, the relationship between star ratings and VADER compound scores is strongly monotonic, validating the consistency of VADER's sentiment scoring across the rating spectrum. Third, word frequency analysis stratified by sentiment category reveals semantically coherent, interpretable vocabulary patterns that provide actionable insights for product management and marketing applications.

The project also identified and characterized systematic misclassification patterns, including VADER's known inability to detect sarcasm and its challenges with long-range negation and domain-specific vocabulary. These limitations motivate clear research directions for future work, including the integration of supervised machine learning classifiers, the exploration of transformer-based contextual language models, and the development of aspect-level sentiment analysis capabilities.

From a broader perspective, this project demonstrates the practical value of NLP and sentiment analysis as tools for converting the massive volume of unstructured consumer review text into structured, actionable intelligence. In an era in which e-commerce platforms generate hundreds of millions of customer reviews annually, automated sentiment analysis systems are not merely academically interesting but operationally indispensable. The framework developed in this project represents a well-validated, interpretable, and extensible foundation for further independent research, personal

portfolio development, and practical applications in the domain of consumer sentiment analysis.

REFERENCES

- [1] B. Pang, L. Lee, and S. Vaithyanathan, "Thumbs up? Sentiment classification using machine learning techniques," in Proceedings of the 2002 Conference on Empirical Methods in Natural Language Processing (EMNLP), pp. 79–86, 2002.
- [2] P. D. Turney, "Thumbs up or thumbs down? Semantic orientation applied to unsupervised classification of reviews," in Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics (ACL), pp. 417–424, 2002.
- [3] C. J. Hutto and E. E. Gilbert, "VADER: A parsimonious rule-based model for sentiment analysis of social media text," in Proceedings of the Eighth International AAAI Conference on Weblogs and Social Media (ICWSM), 2014.
- [4] J. McAuley and J. Leskovec, "From amateurs to connoisseurs: Modeling the evolution of user expertise through online reviews," in Proceedings of the 22nd International World Wide Web Conference (WWW), pp. 897–908, 2013.
- [5] R. Socher, A. Perelygin, J. Wu, J. Chuang, C. D. Manning, A. Ng, and C. Potts, "Recursive deep models for semantic compositionality over a sentiment treebank," in Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing (EMNLP), pp. 1631–1642, 2013.
- [6] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, "BERT: Pre-training of deep bidirectional transformers for language understanding," in Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL-HLT), pp. 4171–4186, 2019.
- [7] A. L. Maas, R. E. Daly, P. T. Pham, D. Huang, A. Y. Ng, and C. Potts, "Learning word vectors for sentiment analysis," in Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics (ACL), pp. 142–150, 2011.
- [8] B. Liu, Sentiment Analysis and Opinion Mining. Morgan & Claypool Publishers, 2012.
- [9] Y. Kim, "Convolutional neural networks for sentence classification," in Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP), pp. 1746–1751, 2014.

- [10] T. Mikolov, K. Chen, G. Corrado, and J. Dean, "Efficient estimation of word representations in vector space," in Proceedings of the International Conference on Learning Representations (ICLR), 2013.
- [11] J. Pennington, R. Socher, and C. Manning, "GloVe: Global vectors for word representation," in Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP), pp. 1532–1543, 2014.
- [12] D. Tang, B. Qin, X. Feng, and T. Liu, "Effective LSTMs for target-dependent sentiment classification," in Proceedings of COLING 2016, pp. 3298–3307, 2016.
- [13] S. Bird, E. Klein, and E. Loper, Natural Language Processing with Python. O'Reilly Media, 2009. [Online]. Available: <https://www.nltk.org/book/>
- [14] W. McKinney, "Data structures for statistical computing in Python," in Proceedings of the 9th Python in Science Conference, pp. 51–56, 2010.
- [15] F. Pedregosa et al., "Scikit-learn: Machine learning in Python," Journal of Machine Learning Research, vol. 12, pp. 2825–2830, 2011.

APPENDIX A: CODE SNIPPETS AND SAMPLE OUTPUTS

A.1 Complete Import and Setup Code

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import re
import string
from collections import Counter
import nltk
from nltk.corpus import stopwords
from nltk.sentiment.vader import SentimentIntensityAnalyzer
from wordcloud import WordCloud

nltk.download('vader_lexicon', quiet=True)
nltk.download('stopwords', quiet=True)
plt.rcParams.update({'figure.dpi': 120})
sns.set_style('whitegrid')
```

A.2 Data Loading and Profiling Code

```
df = pd.read_csv('Reviews.csv')
print(f'Dataset loaded: {df.shape[0]:,} rows, {df.shape[1]} columns')
print(df.isnull().sum())
print(df['Score'].value_counts().sort_index())
```

A.3 Text Preprocessing Function

```
stop_words = set(stopwords.words('english'))

def preprocess_text(text):
    text = text.lower()
    text = re.sub(r'<.*?>', '', text)
    text = re.sub(r'[^a-z\s]', '', text)
    tokens = text.split()
    tokens = [t for t in tokens if t not in stop_words and len(t) > 2]
    return ' '.join(tokens)

df_sample['cleaned_text'] = df_sample['Text'].apply(preprocess_text)
```

A.4 VADER Sentiment Scoring Code

```
sia = SentimentIntensityAnalyzer()

def get_vader_scores(text):
    scores = sia.polarity_scores(str(text))
    return scores['neg'], scores['neu'], scores['pos'],
    scores['compound']

df_sample[['vader_neg', 'vader_neu', 'vader_pos', 'vader_compound']] = \
    df_sample['Text'].apply(lambda x: pd.Series(get_vader_scores(x)))

def label_sentiment(compound):
    if compound >= 0.05: return 'Positive'
    elif compound <= -0.05: return 'Negative'
    else: return 'Neutral'

df_sample['sentiment'] =
df_sample['vader_compound'].apply(label_sentiment)
```

A.5 Agreement Rate Calculation

```
def expected_sentiment(score):
    if score >= 4: return 'Positive'
    elif score <= 2: return 'Negative'
    else: return 'Neutral'

df_sample['expected_sentiment'] =
df_sample['Score'].apply(expected_sentiment)

match = (df_sample['sentiment'] ==
df_sample['expected_sentiment']).mean()

print(f'VADER vs Star Rating agreement: {match:.1%}')
```

A.6 Sample Prediction Output

Table A.1: Sample VADER Scoring Output for Selected Reviews

Score	VADER Compound	VADER Label	Review Excerpt (truncated)
5	0.923	Positive	Absolutely love this coffee — best I have ever tasted, smooth and rich...

Score	VADER Compound	VADER Label	Review Excerpt (truncated)
1	-0.816	Negative	Terrible product. Arrived stale and completely inedible. Do not buy...
3	0.041	Neutral	It was okay. Not quite what I expected but did the job for the price...
1	0.714	Positive	Oh fantastic — yet another product that falls apart immediately! Great job...
4	0.658	Positive	Really good quality, fast delivery. Would definitely recommend to friends...